

# Logic programming with Prolog

Dr. Constantinos Constantinides, P.Eng.


Department of Computer Science and Software Engineering  
Concordia University

[constantinos.constantinides@concordia.ca](mailto:constantinos.constantinides@concordia.ca)

# Programming in logic

- Predicate logic can be used to represent and reason about knowledge.
- We will adopt the Prolog programming language to model and process clauses.
- In this discussion we will use a running example to express the meaning and constraints of data as well as to construct queries over their representation in order to obtain information.

# A running example: A family genealogy tree



# Programs and statements

parent(tom, adam).  
parent(tom, helen).  
parent(sandra, adam).  
parent(sandra, helen).  
parent(michael, andrew).  
parent(michael, john).  
parent(eve, andrew).  
parent(eve, john).  
parent(helen, mark).  
parent(andrew, mark).  
parent(judy, roger).  
parent(judy, jim).  
parent(judy, janis).  
parent(mark, roger).  
parent(mark, jim).  
parent(mark, janis).  
parent(janis, daphne).  
parent(peter, daphne).

Prolog programs consist of collections of statements (called assertions, or clauses).

# Statements and procedures

parent(tom, adam).  
parent(tom, helen).  
parent(sandra, adam).  
parent(sandra, helen).  
parent(michael, andrew).  
parent(michael, john).  
parent(eve, andrew).  
parent(eve, john).  
parent(helen, mark).  
parent(andrew, mark).  
parent(judy, roger).  
parent(judy, jim).  
parent(judy, janis).  
parent(mark, roger).  
parent(mark, jim).  
parent(mark, janis).  
parent(janis, daphne).  
parent(peter, daphne).

Statements are grouped into procedures.

In the example, we have one procedure named **parent**, made up of several statements.

# Procedures and arguments

parent(tom, adam).  
parent(tom, helen).  
parent(sandra, adam).  
parent(sandra, helen).  
parent(michael, andrew).  
parent(michael, john).  
parent(eve, andrew).  
parent(eve, john).  
parent(helen, mark).  
parent(andrew, mark).  
parent(judy, roger).  
parent(judy, jim).  
parent(judy, janis).  
parent(mark, roger).  
parent(mark, jim).  
parent(mark, janis).  
parent(janis, daphne).  
parent(peter, daphne).

Each procedure defines a certain relationship between its arguments.

The programmer decides on how to interpret this relationship.

Here, parent(tom, adam) will be interpreted as “*Tom is the parent of Adam.*”

# Statements revisited: Facts and rules

parent(tom, adam).  
parent(tom, helen).  
parent(sandra, adam).  
parent(sandra, helen).  
parent(michael, andrew).  
parent(michael, john).  
parent(eve, andrew).  
parent(eve, john).  
parent(helen, mark).  
parent(andrew, mark).  
parent(judy, roger).  
parent(judy, jim).  
parent(judy, janis).  
parent(mark, roger).  
parent(mark, jim).  
parent(mark, janis).  
parent(janis, daphne).  
parent(peter, daphne).

There are two kinds of statements:

1) facts, and

2) rules.

Facts are propositions declared to be true.

In the example, the procedure named **parent** consists only by facts.

# Questions and queries

parent(tom, adam).  
parent(tom, helen).  
parent(sandra, adam).  
parent(sandra, helen).  
parent(michael, andrew).  
parent(michael, john).  
parent(eve, andrew).  
parent(eve, john).  
parent(helen, mark).  
parent(andrew, mark).  
parent(judy, roger).  
parent(judy, jim).  
parent(judy, janis).  
parent(mark, roger).  
parent(mark, jim).  
parent(mark, janis).  
parent(janis, daphne).  
parent(peter, daphne).

The collection of statements constitutes a (declarative) database.

We can pose queries on this database.

A query is the codification of a question.

There are only two types of queries:

*1. Is it indeed the case that a given statement is true? (ground query)*

*2. Under what conditions, if any, is a given statement true? (non-ground query)*



# Ground queries

parent(tom, adam).  
parent(tom, helen).  
parent(sandra, adam).  
parent(sandra, helen).  
parent(michael, andrew).  
parent(michael, john).  
parent(eve, andrew).  
parent(eve, john).  
parent(helen, mark).  
parent(andrew, mark).  
parent(judy, roger).  
parent(judy, jim).  
parent(judy, janis).  
parent(mark, roger).  
parent(mark, jim).  
parent(mark, janis).  
parent(janis, daphne).  
parent(peter, daphne).

Ground queries result in a Yes/No (or True/False) response:

For example, the question

*“Is it indeed the case that Peter is the parent of Daphne?”*

will be codified into a query and executed as

**?- parent(peter, daphne).**

# Evaluation of ground queries

parent(tom, adam).  
parent(tom, helen).  
parent(sandra, adam).  
parent(sandra, helen).  
parent(michael, andrew).  
parent(michael, john).  
parent(eve, andrew).  
parent(eve, john).  
parent(helen, mark).  
parent(andrew, mark).  
parent(judy, roger).  
parent(judy, jim).  
parent(judy, janis).  
parent(mark, roger).  
parent(mark, jim).  
parent(mark, janis).  
parent(janis, daphne).  
parent(peter, daphne).

Prolog will take the query

**?- parent(peter, daphne).**

and will start searching the database from top to bottom, one statement at a time trying to match (“unify”) it with a statement.

In trying the first statement, a match (unification) is *not* successful.

# Evaluation of ground queries /cont.

parent(tom, adam).  
parent(tom, helen).  
parent(sandra, adam).  
parent(sandra, helen).  
parent(michael, andrew).  
parent(michael, john).  
parent(eve, andrew).  
parent(eve, john).  
parent(helen, mark).  
parent(andrew, mark).  
parent(judy, roger).  
parent(judy, jim).  
parent(judy, janis).  
parent(mark, roger).  
parent(mark, jim).  
parent(mark, janis).  
parent(janis, daphne).  
parent(peter, daphne).

Prolog will then try to match the query

**?- parent(peter, daphne).**

against the next statement in the program.

Again, if not successful, it will try the next statement in the program...etc. until either a match is found or until the database is exhausted.

# Evaluation of ground queries /cont.

parent(tom, adam).  
parent(tom, helen).  
parent(sandra, adam).  
parent(sandra, helen).  
parent(michael, andrew).  
parent(michael, john).  
parent(eve, andrew).  
parent(eve, john).  
parent(helen, mark).  
parent(andrew, mark).  
parent(judy, roger).  
parent(judy, jim).  
parent(judy, janis).  
parent(mark, roger).  
parent(mark, jim).  
parent(mark, janis).  
parent(janis, daphne).  
parent(peter, daphne). ←

In this example, Prolog will eventually succeed, having matched the query

`?- parent(peter, daphne).`

with the last statement.

Prolog will respond *Yes* (or *true*) to the query.

We have managed to prove that it is indeed the case that `parent(peter, daphne)` is true.

# Non-ground queries

parent(tom, adam).  
parent(tom, helen).  
parent(sandra, adam).  
parent(sandra, helen).  
parent(michael, andrew).  
parent(michael, john).  
parent(eve, andrew).  
parent(eve, john).  
parent(helen, mark).  
parent(andrew, mark).  
parent(judy, roger).  
parent(judy, jim).  
parent(judy, janis).  
parent(mark, roger).  
parent(mark, jim).  
parent(mark, janis).  
parent(janis, daphne).  
parent(peter, daphne).

A non-ground query involves one or more variables.

The question “*Who is a parent of Daphne*” can be transformed into a non-ground query as

**?- parent(X, daphne).**

where **X** (note the capitalization) is a variable.

We are asking Prolog to seek instantiation(s) for variable **X**, provided any exist, that could make the query succeed.


# Unification revisited

Unification (or matching) is a basic operation on terms.

A ground query can unify with a statement, e.g.  
→ **?- parent(tom, adam).**

A non-ground query can unify with a statement only if substitution can be made for any variables so that the two terms can be made equal. In this example, we have

→ **?- parent(X, daphne).**




Tree representation of  
a statement.

Tree representation of a  
non-ground query.

parent(tom, adam).  
parent(tom, helen).  
parent(sandra, adam).  
parent(sandra, helen).  
parent(michael, andrew).  
parent(michael, john).  
parent(eve, andrew).  
parent(eve, john).  
parent(helen, mark).  
parent(andrew, mark).  
parent(judy, roger).  
parent(judy, jim).  
parent(judy, janis).  
parent(mark, roger).  
parent(mark, jim).  
parent(mark, janis).  
parent(janis, daphne).  
parent(peter, daphne).

# Unification revisited /cont.



- The terms `parent(janis, daphne)` and `parent(X, daphne)` unify, instantiating `X` to `janis`.
- There is in fact one more solution, because there are two possible choices to unify with `parent(X, daphne)`.

# Rules: Head and body

A rule is an implication between propositions.

The general form is

`head :- body.`

which reads

*“The head (of the rule) is true, if the body is true.”,*

or, alternatively:

*“The head of the rule can succeed if the body of the rule can succeed.”*



# Formulae and rules

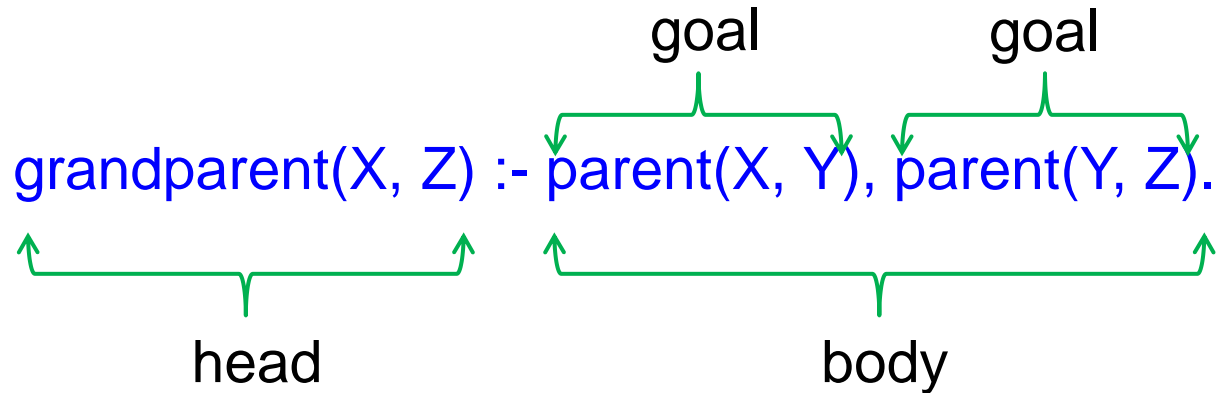
- Let us extend the database with a new procedure `grandparent`.
- Let `p` stand for the binary relation `isParentOf` and let `g` stand for the binary relation `isGrandParentOf`.
- We can define `g` in terms of `p` by the following formula: For persons `x`, `y`, `z`:

$$G = \forall x, y, z ((p(x, z) \wedge p(z, y)) \rightarrow g(x, y))$$

- We can represent the formula with the rule below:  
`grandparent(X, Z) :- parent(X, Y), parent(Y, Z).`

# Rules and goals

Consider the following rule :



The body of rule `grandparent/2` contains two goals.

The goals are related by a conjunction (denoted by the comma symbol).

The rule is now added to the database.

# Evaluation of a ground query in the presence of rules

## An example

parent(tom, adam).  
parent(tom, helen).  
parent(sandra, adam).  
parent(sandra, helen).  
parent(michael, andrew).  
parent(michael, john).  
parent(eve, andrew).  
parent(eve, john).  
parent(helen, mark).  
parent(andrew, mark).  
parent(judy, roger).  
parent(judy, jim).  
parent(judy, janis).  
parent(mark, roger).  
parent(mark, jim).  
parent(mark, janis).  
parent(janis, daphne).  
parent(peter, daphne).

Consider the query `grandparent(judy, daphne)`.

Prolog will search its database from top to bottom, and

a) **unify** the query with the head of the rule,

b) **instantiate** `X` to `judy` and `Z` to `daphne`,

`grandparent(X, Z) :- parent(X, Y), parent(Y, Z).`



`?- grandparent(judy, daphne).`

`grandparent(X, Z) :- parent(X, Y),  
parent(Y, Z).`

# Evaluation of a ground query in the presence of rules

## An example /cont.

parent(tom, adam).  
parent(tom, helen).  
parent(sandra, adam).  
parent(sandra, helen).  
parent(michael, andrew).  
parent(michael, john).  
parent(eve, andrew).  
parent(eve, john).  
parent(helen, mark).  
parent(andrew, mark).  
parent(judy, roger).  
parent(judy, jim).  
parent(judy, janis).  
parent(mark, roger).  
parent(mark, jim).  
parent(mark, janis).  
parent(janis, daphne).  
parent(peter, daphne).  
grandparent(X, Z) :- parent(X, Y),  
                    parent(Y, Z).

c) **resolve** to two new queries (that correspond to the two goals of the rule):

`parent(judy, Y), parent(Y, daphne).`

Both queries must now be evaluated (in the order specified) and if both prove true, then the rule succeeds (and the answer to the query is Yes/True).

# Evaluation of a ground query in the presence of rules

An example /cont.

```
parent(tom, adam).  
parent(tom, helen).  
parent(sandra, adam).  
parent(sandra, helen).  
parent(michael, andrew).  
parent(michael, john).  
parent(eve, andrew).  
parent(eve, john).  
parent(helen, mark).  
parent(andrew, mark).  
parent(judy, roger).  
parent(judy, jim).  
parent(judy, janis).  
parent(mark, roger).  
parent(mark, jim).  
parent(mark, janis).  
parent(janis, daphne).  
parent(peter, daphne).  
grandparent(X, Z) :- parent(X, Y),  
                        parent(Y, Z).
```

The two goals

**parent(judy, Y), parent(Y, daphne).**

will be evaluated as follows:

The first goal **parent(judy, Y)**, will be executed as any other query, unifying with **parent(judy, roger)**, and instantiating **Y** to **roger**.

Once the first goal succeeds, Prolog will try the next one on the right, for the same instantiation:

Can **roger** make the second goal succeed?

No. The query **parent(roger, daphne)** is not successful.

# Evaluation of a ground query in the presence of rules

## An example /cont.

parent(tom, adam).  
parent(tom, helen).  
parent(sandra, adam).  
parent(sandra, helen).  
parent(michael, andrew).  
parent(michael, john).  
parent(eve, andrew).  
parent(eve, john).  
parent(helen, mark).  
parent(andrew, mark).  
parent(judy, roger).  
**parent(judy, jim).**  
parent(judy, janis).  
parent(mark, roger).  
parent(mark, jim).  
parent(mark, janis).  
parent(janis, daphne).  
parent(peter, daphne).  
grandparent(X, Z) :- parent(X, Y),  
                    parent(Y, Z).

Prolog will continue searching the database to find matches that can satisfy both goals

**parent(judy, Y), parent(Y, daphne).**

The first goal, **parent(judy, Y)**, can unify with **parent(judy, jim)**, instantiating **Y** to **jim**.

Can **jim** make the second goal succeed?

No. The query **parent(jim, daphne)** is not successful.

# Evaluation of a ground query in the presence of rules

## An example /cont.

parent(tom, adam).  
parent(tom, helen).  
parent(sandra, adam).  
parent(sandra, helen).  
parent(michael, andrew).  
parent(michael, john).  
parent(eve, andrew).  
parent(eve, john).  
parent(helen, mark).  
parent(andrew, mark).  
parent(judy, roger).  
parent(judy, jim).  
**parent(judy, janis).**  
parent(mark, roger).  
parent(mark, jim).  
parent(mark, janis).  
**parent(janis, daphne).**  
parent(peter, daphne).  
grandparent(X, Z) :- parent(X, Y),  
                    parent(Y, Z).

Prolog will continue searching the database to find matches that can satisfy both goals

**parent(judy, Y), parent(Y, daphne).**

The first goal, **parent(judy, Y)**, can unify with **parent(judy, janis)**, instantiating **Y** to **janis**.

Can **janis** make the second goal succeed?

Yes. The query **parent(janis, daphne)** is successful.

# Evaluation of a non-ground query in the presence of rules: An example

The question

parent(tom, adam).  
parent(tom, helen).  
parent(sandra, adam).  
parent(sandra, helen).  
parent(michael, andrew).  
parent(michael, john).  
parent(eve, andrew).  
parent(eve, john).  
parent(helen, mark).  
parent(andrew, mark).  
parent(judy, roger).  
parent(judy, jim).  
parent(judy, janis).  
parent(mark, roger).  
parent(mark, jim).  
parent(mark, janis).  
parent(janis, daphne).  
parent(peter, daphne).

*“Who are the grandparents of Daphne?”*

is codified into the query

**?- grandparent(G, daphne).**

This will unify with the head of rule  
**grandparent**, instantiating variable **Z** to  
**daphne** and resolving into two goals:

**parent (G, Y), parent(Y, daphne).**

**grandparent(X, Z) :- parent(X, Y),  
parent(Y, Z).**



# Evaluation of a non-ground query in the presence of rules: An example


parent(tom, adam).  
parent(tom, helen).  
parent(sandra, adam).  
parent(sandra, helen).  
parent(michael, andrew).  
parent(michael, john).  
parent(eve, andrew).  
parent(eve, john).  
parent(helen, mark).  
parent(andrew, mark).  
parent(judy, roger).  
parent(judy, jim).  
parent(judy, janis).  
parent(mark, roger).  
parent(mark, jim).  
parent(mark, janis).  
parent(janis, daphne).  
parent(peter, daphne).  
grandparent(X, Z) :- parent(X, Y),  
                          parent(Y, Z).

?- grandparent(X, daphne).

X = judy ;

X = mark ;

No



Upon finding a match, Prolog will stop here and wait for instructions. The semicolon symbol (;) inquires whether more matches can be found.

A period symbol (.) would indicate our intention to stop the search.

# Multi-line rules: Disjunction

- We can now further extend the database with a rule to define the binary relation “is ancestor of.”
- Suppose we let **p** stand for the binary relation **isParentOf** and let **a** stand for the binary relation **isAncestorOf** relation.
- We can now define **a** in terms of **p** by the following formula we will call **A**:

$$A = \forall x, y (p(x, y) \rightarrow a(x, y))$$

$$A = \forall x, y, z ((p(x, z) \text{ and } a(z, y)) \rightarrow a(x, y))$$

## Multi-line rules: Disjunction /cont.

- In other words, **x** is an ancestor of **y** if either **x** is a parent of **y**, or **x** is a parent of an ancestor of **y**. We can represent this in Prolog with the following multi-line rule:

disjunction { ancestor(X, Y) :- parent(X, Y).  
 ancestor(X, Y) :- parent(X, Z), ancestor(Z, Y).  
 conjunction

- A rule can be placed in more than one lines.
- In this case, there is a disjunction between the two lines, and there is a conjunction between the two goals of the body of the second rule.

# Further extending the database

```
man(tom).
man(michael).
man(adam).
man(andrew).
man(john).
man(mark).
man(roger).
man(jim).
man(peter).
woman(sandra).
woman(eve).
woman(helen).
woman(judy).
woman(janis).
woman(daphne).
father(X, Y) :- man(X),
                parent(X, Y).
mother(X, Y) :- woman(X),
                parent(X, Y).
```

We extend the database by introducing four procedures:

**man**, **woman**: made up of facts,

and

**father**, **mother**: made up of rules.

```
father(X, Y)    :- man(X),
                  parent(X, Y).
```

```
mother(X, Y)    :- woman(X),
                  parent(X, Y).
```

# Anonymous variables in rules

- If any parameter of a relation is not important, we can replace it with an anonymous variable denoted by the underscore character (`_`) as follows:

```
is_father(X) :- father(X, _).
```

```
is_mother(X) :- mother(X, _).
```

- We can now pose more queries such as *“Is Tom a father?”*
- To answer this type of question, it does not matter who **Tom** is the father of, as long as **Tom** is found as the first term in a father statement. The query is as follows:

```
?- is_father(tom).
```

```
true
```

# Anonymous variables in queries

- Alternatively we can use anonymous variables in queries, such as

```
?- father(tom, _).
```

```
true
```

# Lists

- Lists are represented in square brackets [...].
- The empty list is represented by [].
- Every non-empty list can be represented in two parts:
  - The head, which is the first element.
  - The tail, which is the list containing the remaining elements.
  - The head of [john, eve, paul] is john.
  - The tail of [john, eve, paul] is [eve, paul].

# List representations

- The symbol `|` in `[H|T]` represents a list whose head is `H` and whose tail is `T`.
- We can represent the above example as `[john | [eve, paul] ]`
- Since `[eve, paul]` is also a list with head `eve` and tail `[paul]`, we can write the above list as `[john | [eve | [paul] ] ]`
- Any one-element list can be written as that element joined to the empty list. Thus, `[paul]` is the same as `[paul | [] ]`
- We can now write the full list as `[john | [eve | [paul | [] ] ] ]`



# Example: Checking for list membership

- We want to define a procedure `member/2` which succeeds if an element is found inside a list.
- We can define list membership recursively as follows:
- Element `X` is a member of a list `L` if  
`X` is the head of `L` (regardless of what the tail is), or  
`member(X, [X|_])`.

or

`X` is a member of the tail of `L` (regardless of what the head is).  
`member(X, [_|T]) :- member(X, T)`.

# Example: Checking for list membership /cont.

- The rule is now

`member(X, [X|_]).`

`member(X, [_|T]) :- member(X, T).`

# Example: Checking for list membership /cont.

## Alternative implementation using the 'cut' operator

- An alternative implementation of `member/2` is shown:

```
member(X, [X|_]) :- !.
```

```
member(X, [_|T]) :- member(X, T).
```

The `!` tells the interpreter to stop looking for alternatives once the LHS succeeds.

However, if `member(X, [X|_])` fails it will check `member(X, [_|T])`.

# Example: Checking for list membership /cont.

## Initial implementation

?- member(a, [a, b, c]).  
true

?- member(a, []).  
false.

?- member([b, c], [[a], [b, c]]).  
true

?- member(X, [a, b]).  
X = a ;  
X = b ;  
false.

## Implementation using 'cut'

?- member(a, [a, b, c]).

?- member(a, []).  
false.

?- member([b, c], [[a], [b, c]]).

?- member(X, [a, b]).  
X = a ;  
false.

# Example: The last element in a list

- In this example, we want to define a rule `last/2` which succeeds if an element is found to be in the last position of a non-empty list.
- We can identify two cases for this:
  1. The list has one element.
  2. The list has more than one element.

# Example: The last element in a list /cont.

- Case 1: The list has only one element.
- In this case, the last element is the only existing element of the list.
- The following rule,

`last(L, [L]).`

reads “*The rule succeeds if element  $L$  is found to be the only element of a given list.*”

# Example: The last element in a list /cont.

- Case 2: The list has more than one element.
- In this case, we need to reduce the problem to the one that can be handled by case 1.
- In other words, the clause will succeed once it chops off all elements, one by one, until it ends up with one element.

# Example: The last element in a list /cont.

- The following rule,

```
last(L, [H|T]) :- last(L, T).
```

reads “*The rule can succeed for a list whose head is  $H$  and whose tail is  $T$ , if it can succeed for a new list which is the tail  $T$  of the original list.*”



## Example: The last element in a list /cont.

- In other words, let us get rid of the first element and see if we end up with only one element in which case the rule of case 1 will determine that this remaining element is indeed the last element.
- However, if after getting rid of the first element we end up with a list with more than one elements, we must repeat this chopping off the head of the list, until we end up with a list which has only one element and subsequently handled by the first rule (of case 1).

# Example: The last element in a list /cont.

## Evaluation of a ground query [1 of 3].

Given the rule,

```
last(L, [L]).
```

```
last(L, [H|T]) :- last(L, T).
```

Consider the query

```
?- last(c, [a, b, c]).
```

Prolog will search its database (from top to bottom) and

**unify** the query with the head of the second statement of the rule,

**instantiate** variable **L** to **c** and variable **[H|T]** to **[a | b, c]**,

**resolve** to a new query (that corresponds to the body of the rule):

```
last(c, [b, c]).
```

This new query must now be evaluated.

# Example: The last element in a list /cont.

## Evaluation of a ground query [2 of 3].

Given the rule,

```
last(L, [L]).
```

```
last(L, [H|T]) :- last(L, T).
```

The query

```
?- last(c, [b, c]).
```

Will cause Prolog to (perform a new search and)

**unify** the query with the head of the second statement of the rule,

**instantiate** variable `L` to `c` and variable `[H|T]` to `[b | c]`,

**resolve** to a new query (that corresponds to the body of the rule):

```
last(c, [c]).
```

This new query must now be evaluated.

# Example: The last element in a list /cont.

## Evaluation of a ground query [3 of 3].

Given the rule

```
last(L, [L]).  
last(L, [H|T]) :- last(L, T).
```

The query

```
?- last(c, [c]).
```

Will cause Prolog to (perform a new search and)

**unify** the query with the head of the first statement of the rule, and yield a success, indicating that it is indeed the case that **c** is found in the last position of the list.

# Arithmetic, relational and logical operators

- Available arithmetic operators:  $+$ ,  $-$ ,  $*$ ,  $/$ ,  $**$ ,  $\text{mod}$ .
- Keyword  $\text{is}$  denotes an assignment.
- Relational operators:  $<$ ,  $>$ ,  $=<$ ,  $>=$ ,  $:=$ ,  $=\backslash=$ 
  - $=$  is used for unification, e.g.  $X = 7$
  - $==$  is used to denote identical terms, e.g.  $5 == 5$
  - $\text{is}$  denotes an assignment and evaluates the RHS, e.g.  $X \text{ is } (1 + 2)$
  - $:=$  evaluates both sides, e.g.  $(1 + 3) := (2 + 2)$
- Logical operators:  $\backslash+$

# Example: Find the largest between two numbers

## Correct (using 'cut')

```
larger(X, Y, X) :- X > Y, !.  
larger(X, Y, Y).
```

```
?- larger(5, 7, X).  
X = 7 .
```

```
?- larger(11, 7, X).  
X = 11 ;  
false.
```

## Incorrect

```
larger(X, Y, X) :- X > Y.  
larger(X, Y, Y).
```

```
?- larger(11, 7, X).  
X = 11 ;  
X = 7 ;  
false.
```

# Example: Calculating the size of a list

- We can recursively obtain the length of a list as follows:
- If the list is empty, then its length is 0.
- If the list is not empty, then its length is 1 + length of its tail.
- Consider rule `size/2` to read in a list and calculate its length.

```
size([],0).
```

```
size([H|T],N) :- size(T,N1), N is N1+1.
```

# Example: Calculating the size of a list

- We can execute queries as follows:

```
?- size([],N).
```

```
N = 0.
```

```
?- size([a,b,c],N).
```

```
N = 3.
```

```
?- size([[a,b],c],N).
```

```
N = 2.
```

```
?- size([[a,b,c]]),N).
```

```
N = 1.
```



# Example: Calculating the size of a list

- We can execute queries as follows:

```
?- delete(a, [a], NewList).
```

```
NewList = [] .
```

```
?- delete(a, [a, e, o], NewList).
```

```
NewList = [e, o] .
```

```
?- delete(a, [e, o, a], NewList).
```

```
NewList = [e, o] .
```

# Example: Deleting an element from a list

- Consider a list **L** and an element **X** that we need to delete from **L**.
- We can identify three cases:
  1. If **X** is the only element, then the result is the empty list.
  2. If **X** is head of **L**, then the result is the tail of the initial list.
  3. If **X** is present in the tail part, then keep the head and recur on the tail.
- Consider rule **delete/3**:

**delete(X, [X], []).**

**delete(X, [X|L1], L1).**

**delete(X, [Y|L2], [Y|L1]) :- delete(X, L2, L1).**

# Built-in utility functions

- The built-in function `findall(X, P, L)` returns a list `L` with all values for `X` that satisfy predicate `P`.
- To eliminate redundancies in a list, we can use the built-in function `list_to_set(List, Set)` that converts the list (with possibly repeated elements) into a set.
- The built-in function `length(List, L)` returns the length `L` of a given list.

# Example with **findall** and **list\_to\_set** in a query

- Let us obtain a set of all fathers:

```
?- forall(F, father(F, _), Lst).
```

```
Lst = [tom, tom, michael, michael, andrew, mark,  
       mark, mark, peter].
```

```
?- forall(F, father(F, _), Lst), list_to_set(Lst, Set).
```

```
Lst = [tom, tom, michael, michael, andrew, mark,  
       mark, mark, peter],  
Set = [tom, michael, andrew, mark, peter].
```

# Example with **findall** and **list\_to\_set** in a rule

- The query  
`findall(F, father(F, _), Lst), list_to_set(Lst, Set)`  
is rather long and complex.
- We can encapsulate its size and complexity in a rule  
`get_all_fathers/1`:

```
get_all_fathers(Set) :-  
    forall(F, father(F, _), Lst),  
    list_to_set(Lst, Set).
```

```
?- get_all_fathers(Set).
```

```
Set = [tom, michael, andrew, mark, peter].
```

# Example with **findall** and **length** in a rule

- Let us construct a rule `qualifies_for_benefits/1` that succeeds if the argument is a mother of at least three children.

```
qualifies_for_benefits(P) :-  
    woman(P),  
    findall(P, parent(P, _), L),  
    length(L, N),  
    N >= 3.
```

```
?- qualifies_for_benefits(Name).
```

```
Name = judy ;
```

```
false.
```